

PFLOCK Report

Andres Calderon

University of California, Riverside

February 7, 2020


Experiment setup...

1. Running a distance self-join to find the pairs of points far each other less than ε .
2. Using LA_25K dataset (time instant 19), quadtree partitioner and quadtree local indexing.
3. $\varepsilon = 20$, *partitions* ≈ 150 (3x number of available cores).
4. 18 executors, 3 cores each.

Partitions performace...

Execution time for partitions during a distance self-join

Showing 6 of 10 runs...



Finding pairs algorithm...

Algorithm 1 FINDPAIRS algorithm

Require: a dataset of points \mathcal{P} , a number of partitions p and a distance threshold ε .

- 1: **function** FINDPAIRS ($\mathcal{P}, p, \varepsilon$)
 - 2: Partition \mathcal{P} using a Quadtree and p partitions ▷ Using Algorithm 1 in next page.
 - 3: Create a circle of radius ε for each point in \mathcal{P} and store them in \mathcal{Q} ▷ keep same id
 - 4: Partition \mathcal{Q} using the same partitioner of \mathcal{P}
 - 5: Build local index in \mathcal{P} ▷ Using operations provided by GeoSpark
 - 6: Execute a distance join query in \mathcal{P} and \mathcal{Q} using ε as distance ▷ Using Algorithm 4 in following pages.
 - 7: Filter those pairs where $p_1.id < p_2.id$
 - 8: **end function**
-

GeoSpark Partitioning algorithm...

Algorithm 1 SRDD spatial partitioning

Data: An original SRDD

Result: A repartitioned SRDD

/ Step 1: Build a global grid file at master node */*

- 1 Take samples from the original SRDD *A* partitions in parallel;
- 2 Construct the selected spatial structure on the collected sample at master node;
- 3 Retrieve the grids from built spatial structures;

/ Step 2: Assign grid ID to each object in parallel */*

4 **foreach** *spatial object in SRDD A* **do**

5 **foreach** *grid* **do**

6 **if** *the grid intersects the object* **then**

7 Add (grid ID, object) pair into SRDD *B*;

// Only needed for R-Tree partitioning

8 **if** *no grid intersects the object* **then**

9 Add (overflow grid ID, object) pair into SRDD *B*;

/ Step 3: Repartition SRDD across the cluster */*

10 Partition SRDD *B* by ID and get SRDD *C*;

11 Cache the new SRDD *C* in memory and return it;

GeoSpark GSJoin algorithm...

Algorithm 4 *GSJoin* algorithm for range join and distance join query

Data: (repartitioned) SRDD A and (repartitioned) SRDD B

Result: PairRDD in schema <Left object from A, right object from B>

/ Step1: Zip partitions */*

1 **foreach** *partition pair from SRDD A and B with the same grid ID i* **do**

2 | Merge two partitions to a bigger partition that has two sub-partitions;

3 Return the intermediate SRDD C;

/ Step2: Run partition-level local join */*

4 **foreach** *partition P in the C* **do**

5 | **foreach** *object O_A in the sub-partition from A* **do**

6 | | **if** *an index exists in the sub-partition from B* **then**

| | | // Filter phase

7 | | | Query the spatial index of this partition using the O_A 's MBR;

| | | // Refine phase

8 | | | Check the spatial relation using real shapes of O_A and candidate objects O_B s;

| | | */* Step3: Remove duplicates */*

9 | | | Report $\langle O_A, O_B \rangle$ pair only if the reference point of this pair is in P;

10 | | **else**

11 | | | **foreach** *object O_B in the sub-partition from B* **do**

12 | | | | Check spatial relation between O_A and O_B ;

| | | | */* Step3: Remove duplicates */*

13 | | | | Report $\langle O_A, O_B \rangle$ pair only if the reference point of this pair is in P;

14 Generate the result PairRDD;
